



**Politecnico
di Torino**



Advanced Machine Learning

A.A. 2023/2024

Tatiana Tommasi



**Politecnico
di Torino**



A short survey

Do you know what is Principal Component Analysis (PCA)?

Do you know what is Support Vector Machines (SVM)?

Have you ever trained an SVM classification model?

Do you know what is cross validation?

Do you have experience with python coding?

Have you ever trained a CNN classification model?

Basic Information

Teacher: Tatiana Tommasi

<http://www.tatianatommasi.com/>
tatiana.tommasi@polito.it

Support from:

Antonio Alliegro, Paolo Rabino, Leonardo Iurada, Claudia Cuttano...

Prof. Barbara Caputo, Giuseppe Averta, Carlo Masone, Raffaello Camoriano

Visual and Multimodal Applied Learning Laboratory (Vandal Lab)

<http://vandal.polito.it/>

Code: 01URWOV

Master of science in Computer Engineering (Ingegneria Informatica)

Credits: 6CFU

Basic Information

When/Where: Tuesday 16:00-19:00, room 3s
Friday 11:30-13:00, room 1s

How: in presence + online (rare pre-recorded videos) + discussion sessions
Exercises: in presence
everything will be recorded with exception of the exercise sessions.

Slack channel:
<https://join.slack.com/t/aml2023/signup>
you should access with your @studenti.polito.it email account
It will be useful mainly to interact with the project mentors

Useful Material

Deep Learning. Ian Goodfellow, Yoshua Bengio, Aaron Courville, <https://www.deeplearningbook.org/>

The little book of Deep Learning <https://fleuret.org/francois/lbdl.html>


Artificial Intelligence A Modern Approach. Stuart J. Russel and Peter Norvig

An Introduction to Probability Theory and its applications. William Feller

Machine Learning: a probabilistic perspective. K. P. Murphy.

Understanding Machine Learning: From Theory to Algorithms. Shai Shalev-Shwartz and Shai Ben-David

CVPR
NeurIPS
ECCV
ICCV
ICLR
...
arXiv



Useful Material

Deep Learning. Ian Goodfellow, Yoshua Bengio, Aaron Courville, <https://www.deeplearningbook.org/>

The little book of Deep Learning <https://fleuret.org/francois/lbdl.html>



Artificial Intelligence A Modern Approach. Stuart J. Russel and Peter Norvig

An Introduction to Probability Theory and its applications. William Feller

Machine Learning: a probabilistic perspective. K. P. Murphy.

Understanding Machine Learning: From Theory to Algorithms. Shai Shalev-Shwartz and Shai Ben-David

CVPR
NeurIPS
ECCV
ICCV
ICLR
...
arXiv



Useful Material

Deep Learning. Ian Goodfellow, Yoshua Bengio, Aaron Courville, <https://www.deeplearningbook.org/>

The little book of Deep Learning <https://fleuret.org/francois/lbdl.html>

Artificial Intelligence A Modern Approach. Stuart J. Russel and Peter Norvig

An Introduction to Probability Theory and its applications. William Feller

Machine Learning: a probabilistic perspective. K. P. Murphy.

Understanding Machine Learning: From Theory to Algorithms. Shai Shalev-Shwartz and Shai Ben-David

Specific references will be given throughout the course in the form of **book chapters** and **scientific articles**.

Courses:

online - Convolutional Neural Networks for Visual Recognition - Fei Fei Li - <http://cs231n.stanford.edu/>

online - Machine Learning course on Coursera - Andrew Ng - <https://www.coursera.org/>

polito - Data Spaces - 01RLPOV, 01RLPNG, Prof. Vaccarino

polito - Machine Learning and Pattern Recognition - 01URTOV, Prof. Cumani

Pre-requirement: Python basic elements

Online tutorial for Python: <https://docs.python.org/3/tutorial/>

Online tutorial for Pytorch: <https://pytorch.org/tutorials/>

Topics and Objective

- AI, ML basic definitions
- basic probability
- loss, risk, prediction rules
- ~~shallow supervised/unsupervised learning~~
- deep learning (supervised, unsupervised, recurrent, style transfer)
- **computer vision applications**

Topics and Objective

- AI, ML basic definitions
- basic probability
- loss, risk, prediction rules
- ~~shallow supervised/unsupervised learning~~
- deep learning (supervised, unsupervised, recurrent, style transfer)
- **computer vision applications**

+ 3d geometry

+ natural language processing

Topics and Objective

- AI, ML basic definitions
- basic probability
- loss, risk, prediction rules
- ~~shallow supervised/unsupervised learning~~
- deep learning (supervised, unsupervised, recurrent, style transfer)
- **computer vision applications**

+ 3d geometry

+ natural language processing



e l l i s

European Laboratory for Learning and Intelligent Systems

On the Map

List of units

Alicante	Amsterdam	Berlin	Cambridge
Copenhagen	Darmstadt	Delft	Edinburgh
Lausanne (EPFL)	Zürich (ETH)	Freiburg	Genoa
Haifa (Technion)	Heidelberg	Helsinki	Jena
Vienna (IST Austria)	Leuven	Linz	Lisbon
London (UCL)	Madrid	Manchester	Milan
Modena (Unimore)	Munich	Nijmegen	Oxford
Paris	Prague	Saarbrücken	Stuttgart
Tel Aviv	Tübingen	Turin	

Unit Director



Tatiana Tommasi
Scholar

Unit Members



Barbara Caputo
Fellow



Diego Valsesia
Member




Enrico Magli
Fellow



Giulia Fracastoro
Member



Giuseppe Averta
Member



Giuseppe Rizzo
Member



Raffaello Camoriano
Member

On the Map

List of units

Alicante	Amsterdam	Berlin	Cambridge
Copenhagen	Darmstadt	Delft	Edinburgh
Lausanne (EPFL)	Zürich (ETH)	Freiburg	Genoa
Haifa (Technion)	Heidelberg	Helsinki	Jena
Vienna (IST Austria)	Leuven	Linz	Lisbon
London (UCL)	Madrid	Manchester	Milan
Modena (Unimore)	Munich	Nijmegen	Oxford
Paris	Prague	Saarbrücken	Stuttgart
Tel Aviv	Tübingen	Turin	

Unit Director



Tatiana Tommasi
Scholar

Unit Members



Barbara Caputo
Fellow



Diego Valsesia
Member




Enrico Magli
Fellow



Giulia Fracastoro
Member



Giuseppe Averta
Member



Giuseppe Rizzo
Member



Raffaello Camoriano
Member



Google
Scholar



Course Aim

- You will get to know the basics of ML & DL
- You will learn how to read & understand a paper (code included)
- You will learn how to design & implement your own ML/DL algorithm for a given problem
- You will learn how to explain what you know about ML/DL to others

Course Aim

- You will get to know the basics of ML & DL

~~mix of in-person, online live and recorded lectures~~

- You will learn how to read & understand a paper (code included)

individual/small group interactive work guided by me and by teaching assistants

- You will learn how to design & implement your own ML/DL algorithm for a given problem

work in groups mainly guided by the teaching assistants

- You will learn how to explain what you know about ML/DL to others

presentation in front of the class / at the exam

Exam Modality

Paper Presentation + Project + Oral Exam

Exam Modality

Paper Presentation + Project + Oral Exam

By mid-November (it might vary a bit) we will have covered basics DL

Then the lectures will be dedicated to presentation of seminal papers of the last 3-7 years (roughly) on various DL topics. The topics will be of interest for all projects.

You will present papers in groups: **3 students, 2 papers**

Papers will be assigned by me in random order.

You will know the title of your paper 2 weeks before your presentation. There will be some time to adjust the calendar.

You will have **10 minutes sharp** for the presentation.

You will be judged on your comprehension of the papers and on your ability to explain it.

Minimum individual score: 6

Maximum individual score : 10

Exam Modality

Paper Presentation + Project + Oral Exam

By mid-November (it might vary a bit) we will have covered basics DL

Then the lectures will be dedicated to presentation of seminal papers of the last 3-7 years (roughly) on various DL topics. The topics will be of interest for all projects.

HOW

There will be a lecture dedicate to how to present a paper.

WHY????

Because working with ML/DL means being a life-long-learner...

Exam Modality

Paper Presentation + Project + Oral Exam

At the beginning of November one lecture will be dedicated to possible project research-areas.

There will be >4 proposals, for each proposal there will be some mandatory work plus suggestions on how to extend it.

You will work in teams of 3 people.

Project Proposal Submission: Before Christmas break (more details on date / submission modality will come), you will indicate: (1) your team (2) the project you chose (3) your planned extension

Final report due: ten days before the oral exam

The report will be a pdf of max 8 pages using a provided template

You can keep your project till the last exam session before the next course beginning (1 year)

Minimum group score: 6

Maximum group score : 10

Exam Modality

Paper Presentation + Project + Oral Exam

You will make a group presentation (about 15 mins) of your project and a Q&A session will follow

Minimum individual score : 6

Maximum individual score : 10

Important Note (1) : This is not a trial-error exam

Important Note (2): the code of your project will be checked - don't cheat

Questions: Paper Presentation

Paper Presentation:

- Can I chose the paper?
No
- I cannot make the presentation on my assigned date because ...
You can ask a colleague for a switch: in this case I have to receive an email in advance explaining why the switch happens and who will present the paper. No impact on the score either way.
- I do not make the presentation on the assigned day & I don't make the switch...
You get 0 points
- Will I get guidance/help in preparing the presentation?
Yes I will make one lecture on how to read a scientific paper & one lecture on how to present a scientific paper/work. You will not get individual help or support to try out the presentation.

Questions: Project

- Can I propose my own project?
No
- Can I do the project alone/in two?
No. If you don't know anyone we'll match you
- What computing facilities I can use for the project?
Google Colab
- How will I chose the project?
There will be >4 projects, each led by a different teaching assistant.
In November each project will be presented with a lecture & live Q&A to help you choosing
- I am not proficient in python, will I be able to make the project?
It will be tough but you can exploit a lot of online material to learn python...if you want to make a living in ML you better learn it!
- Will I get guidance/help in working on the project?
Yes you will be supported by the teaching assistant both with meetings and on slack.

Questions: Oral

- Can I make the oral exam in a different session than the rest of my group?
No
- Is the oral exam in English?
Yes
- What will be asked during the exam?
Everything related to the project, from implementation details to theory related to the algorithms used in the project. There will be also questions on any topic explained during the lectures.
- Why are we doing this project???? 
This is the main question I'm going to ask you during the exam...
- Can I reject the final vote and sit again the exam?
Yes. For the next exam session you will need to search for another team on a different topic (i.e. you will have to work on a different project). Please come to sit at the exam when you are prepared.

Why Are you attending this course?

Learn about basic tools that you will need for:

- **your M. Sc./PhD Thesis**
- **your Start Up**
- **Wall Street**
- **Google, Facebook, Amazon, Snapchat, ...**
- **Huawei, Samsung, IBM, Siemens, Microsoft, Apple,...**
- **Disney, Netflix...**
- **Linkedin, AirBnB, Booking,...**
- **Uber, BMW, Honda, Toyota, Tesla, FCA, iRobot, ...**

MACHINE INTELLIGENCE 3.0

ENTERPRISE INTELLIGENCE

VISUAL

Orbital Insight planet
clarifai DEEP VISION
cortica Ilogicon
SPACE_KNOW Captricity
netra deepomatic

AUDIO

Gridspace TalkIQ
nexidia twilio
CAPIO Expect Labs
Clover Mobvoi
QuirousAI popUP archive

SENSOR

PREDIX GIoT MAANA
Sentenai PLANET OS
UPTAKE IMUBIT Preferred Networks
thingworx KONUX Alluvium

INTERNAL DATA

PRIMER IBM WATSON
Gycorp Palantir ARIMO
Alation Sapho Outlier
Digital Reasoning

MARKET

mattermark Quid
Datafox PREMISE
Bottlenose MOTIVA
enigma CB INSIGHTS
Tracxn predata

CUSTOMER SUPPORT

DigitalGenius Kasisto
Eloquent WISEIO
ACTIONIQ zendesk
Preact CLARABRIDGE

SALES

collective[i] osense
fuse|machines AVISO
salesforce INSIDE SALES
Zenight clari

MARKETING

MINTIGO Lattice RADIUS
LiftIgniter PERSADO
brightfunnel retention
COGNICOR AIRPR msgai

SECURITY

CYCLANCE DARKTRACE
ZIMPERIUM deepinstinct
Sentinel DEMISTO
graphistry drawbridge
SignalSense AppZen

RECRUITING

textio entelo
Wade & Wendy hiQ
unifive SpringRole
GIGSTER HireVue

AUTONOMOUS SYSTEMS

GROUND NAVIGATION

drive.ai AdasWorks
ZOOX MOBILEYE
UBER Google TESLA
nuro Autonomy

AERIAL

SKYDIO SHIELD AI
Airware DJI LILY
DroneDeploy
pilot.ai SKYCATCH

INDUSTRIAL

JAYBRIDGE OSARO
CLEARPATH fetch
KINDRED rethink robotics
HARVEST

PERSONAL

amazon alexa
Cortana Allo
facebook
Siri Replika

AGENTS

PROFESSIONAL

butter.ai pogo SKIPFLAG
clara x.ai slack
talla Zoom sudo

INDUSTRIES

AGRICULTURE

BLUERIVER mavrx
tile TRACE pivot Bio
TerraVision AGRI-DATA
Descartes Labs udie

EDUCATION

KNEWTON volley
gradescope
CTI coursera
UDACITY alt school

INVESTMENT

Bloomberg sentient
ISENTIUM KENSHO
alphasense Dataminr
CEREBELLUM CAPITAL Quandl

LEGAL

blue J BEAGLE
Everlaw RAVEL
Seal ROSS
LEGAL ROBOT

LOGISTICS

NAUTO Acerta
PRETECKT
Routific clearmetal
MARBLE PITSTOP

INDUSTRIES CONT'D

MATERIALS

zymergen Citrine
Eigen Innovations
SIGHT MACHINE
SINKBO nanotronics
BIOWORKS CALCULARIO

RETAIL FINANCE

TALA zest finance
Lendo earnest
affirm MIRADOR
wealthfront Betterment

HEALTHCARE

PATIENT

PULSE CareSkore
ZEPHYRUS HEALTH IBM Watson Health
Oncoda SENTRIAN
Atomwise Numerate

IMAGE

BUTTERFLY 3SCAN
ARTERYS enlitic
BAYLABS imagia
Google DeepMind

BIOLOGICAL

iCarbonX color GRAIL
deep genomics RECURSION
LUMINIST Numerate
Atomwise verily WHOLE BIO

TECHNOLOGY STACK

AGENT ENABLERS

OCTANE.AI howdy. Maluuba KITT.AI
OpenAI Gym Kasisto AUTOMAT
semantic machines

DATA SCIENCE

DOMINO SPARKBEYOND rapidminer
kaggle DataRobot yhat AYASDI
data iku seldon yseop bigml

MACHINE LEARNING

CognitiveScale GoogleML context relevant
Cycorp HyperScience narda logics minds.ai H2O.ai
SCALED INFERENCE spark cognition loop GEOMETRIC INTELLIGENCE
deep sense.io reactive sky mind bonsai

NATURAL LANGUAGE

agolo FLYLIEN LEXALYTICS
Narrative Science loop@ spaCy LUMINOSO
cortical.io MonkeyLearn

DEVELOPMENT

SIGOPT HyperOpt fuzzyio okite
rainforest lobe Anodot
Signifai LAYER6 bonsai

DATA CAPTURE

CrowdFlower Diffbot CrowdAI import io
Paxata DATA SIFT amazon mechanical turk enigma
WorkFusion DATALOGUE TRIFACTA parsehub

OPEN SOURCE LIBRARIES

Keras Chainer CNTK TensorFlow Caffe
H2O DEEPLARNING4J theano torch
DSSTNE Scikit-learn AzureML neon
MXNet DMTK Spark PaddlePaddle WEKA

HARDWARE

KNUPATH TENSTORRENT Cirrascale
NVIDIA intel nervana Movidius
tensilica GoogleTPU 1026 Labs Qualcomm
Cerebras Isosemi

RESEARCH

OpenAI nnsense ELEMENT vicarious
KNOGIN Numenta Kimera Systems Cogital

Why Are you attending this course?

Learn about basic tools that you will need for:

- **your M. Sc./PhD Thesis**
- **your Start Up**
- **Wall Street**
- **Google, Facebook, Amazon, Snapchat, ...**
- **Huawei, Samsung, IBM, Siemens, Microsoft, Apple,...**
- **Disney, Netflix...**
- **Linkedin, AirBnB, Booking,...**
- **Uber, BMW, Honda, Toyota, Tesla, FCA, iRobot, ...**

Caveat...

- NOT: learn to use deep learning as a black box
- YES: understand models and architectures

Why Are you attending this course?

Learn about basic tools that you will need for:

- **your M. Sc./PhD Thesis**
- **your Start Up**
- **Wall Street**
- **Google, Facebook, Amazon, Snapchat, ...**
- **Huawei, Samsung, IBM, Siemens, Microsoft, Apple,...**
- **Disney, Netflix...**
- **Linkedin, AirBnB, Booking,...**
- **Uber, BMW, Honda, Toyota, Tesla, FCA, iRobot, ...**

Caveat...

- NOT: learn to use deep learning as a black box
- YES: understand models and architectures

This is not an easy course, but results speak for themselves!

Why Are you attending this course?

Top 10 jobs involving AI skills

Top jobs seeking machine learning or artificial intelligence skills

Rank	Job title	% of postings containing AI or machine learning	Rank	Job title	% of postings containing AI or machine learning
1.	Machine learning engineer	75.0%	6.	Algorithm developer	46.9%
2.	Deep learning engineer	60.9%	7.	Junior data scientist	45.7%
3.	Senior data scientist	58.1%	8.	Developer consultant	44.5%
4.	Computer vision engineer	55.2%	9.	Director of data science	41.5%
5.	Data scientist	52.1%	10.	Lead data scientist	32.7%

Source: Indeed




AI job titles with the highest salaries




Rank	Job title	Average salary
1.	Machine learning engineer	\$142,858.57
2.	Data scientist	\$126,927.41
3.	Computer vision engineer	\$126,399.81
4.	Data warehouse architect	\$126,008.25
5.	Algorithm engineer	\$109,313.51


What is AI?



What is AI?




What is AI?




Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

What is AI?


Thinking
humanly




Acting
humanly



Thinking
rationally




Acting
rationally



<http://www-inst.eecs.berkeley.edu/~cs188/fa19/>
<http://slazebni.cs.illinois.edu/fall17/>

What is AI?

Thinking
humanly




Acting
humanly



Make computers think like a human
How does a human think?


- brain imaging - neuroscience
- psychological experiments - cognitive science



Alison Gopnik - developmental scientist

What is AI?

Thinking
humanly




Acting
humanly




Make computers think like a human
How does a human think?

- brain imaging - neuroscience
- psychological experiments - cognitive science




Alison Gopnik - developmental scientist

The Turing Test Approach



What is AI?

Thinking
humanly




Acting
humanly




Make computers think like a human
How does a human think?

- brain imaging - neuroscience
- psychological experiments - cognitive science



Alison Gopnik - developmental scientist


The Turing Test Approach



- **natural language processing** - communicate
- **knowledge representation** - store knowledge
- **automated reasoning** - answer questions and draw conclusions
- **machine learning** - detect and extract patterns, adapt to new circumstances
- **computer vision** - perceive objects
- **robotics** - manipulate objects

What is AI?

Thinking
humanly




Acting
humanly




Make computers think like a human
How does a human think?

- brain imaging - neuroscience
- psychological experiments - cognitive science



Alison Gopnik - developmental scientist

The Turing Test Approach



- **natural language processing** - communicate
- **knowledge representation** - store knowledge
- **automated reasoning** - answer questions and draw conclusions
- **machine learning** - detect and extract patterns, adapt to new circumstances
- **computer vision** - perceive objects
- **robotics** - manipulate objects

Aeronautical engineering “*goal of the field is making machines that flies so exactly like pigeons that they can fool other pigeons*”

What is AI?



Acting
humanly

Issues with this test:

- Variability in protocols, judges
- Success depends on deception!
- Chatbots can do well using “cheap tricks” (Eliza 1966)

Winograd schema: Multiple choice questions that can be easily answered by people but cannot be answered by computers using “cheap tricks”


The sack of potatoes had been placed below the bag of flour, so it had to be moved first. What had to be moved first?

- The sack of potatoes
- The bag of flour

<https://ptchallenge-workshop.github.io/>



The Turing Test Approach




- **natural language processing** - communicate
- **knowledge representation** - store knowledge
- **automated reasoning** - answer questions and draw conclusions
- **machine learning** - detect and extract patterns, adapt to new circumstances
- **computer vision** - perceive objects
- **robotics** - manipulate objects

Aeronautical engineering “*goal of the field is making machines that flies so exactly like pigeons that they can fool other pigeons*”

What is AI?

Thinking
rationally



Acting
rationally

Idealized or “right” way of thinking

Logic: patterns of argument that always yield correct conclusions when supplied with correct premises


Logicist approach to AI: describe problem in formal logical notation and apply general deduction procedures to solve it

Issues:

- Describing real-world problems and knowledge in logical notation
- Dealing with uncertainty
- Computational complexity of finding the solution starting from a list of facts

What is AI?

Thinking
rationally



Acting
rationally

Idealized or “right” way of thinking

Logic: patterns of argument that always yield correct conclusions when supplied with correct premises

Logicist approach to AI: describe problem in formal logical notation and apply general deduction procedures to solve it

Issues:


- Describing real-world problems and knowledge in logical notation
- Dealing with uncertainty
- Computational complexity of finding the solution starting from a list of facts

A rational agent acts so as to achieve the best outcome or, when there is uncertainty, the best expected outcome.

- Goals are expressed in terms of the utility of outcomes (quantitative measure)
- Naturally accommodates uncertainty
- Rationality only concerns what decisions are made (no cognitive process behind)
- Being rational means maximizing your expected utility


In practice, utility optimization is subject to the agent’s computational constraints (bounded rationality or bounded optimality)

What is AI?



Acting
rationally

AI = Computational Rationality




A rational agent acts so as to achieve the best outcome or, when there is uncertainty, the best expected outcome.

- Goals are expressed in terms of the utility of outcomes (quantitative measure)
- Naturally accommodates uncertainty
- Rationality only concerns what decisions are made (no cognitive process behind)
- Being rational means maximizing your expected utility

In practice, utility optimization is subject to the agent's computational constraints (bounded rationality or bounded optimality)


<http://www-inst.eecs.berkeley.edu/~cs188/fa19/>

What is AI?



Acting
rationally

AI = Computational Rationality



<http://www-inst.eecs.berkeley.edu/~cs188/fa19/>

All the skills needed for the Turing Test allow an agent to act rationally

- **knowledge representation** - store knowledge
- **automated reasoning** - answer questions and draw conclusions


enable the agent to reach good decisions

- **natural language processing** - communicate
generate comprehensible sentences and get by in a complex society

- **machine learning** - detect and extract patterns, adapt to new circumstances
improve ability to generate effective behavior


- **computer vision** - perceive objects
- **robotics** - manipulate objects
interact beyond language communication

What is AI?



Acting
rationally

AI = Computational Rationality



<http://www-inst.eecs.berkeley.edu/~cs188/fa19/>

All the skills needed for the Turing Test allow an agent to act rationally

- **knowledge representation** - store knowledge
- **automated reasoning** - answer questions and draw conclusions

enable the agent to reach good decisions

- **natural language processing** - communicate
generate comprehensible sentences and get by in a complex society

- **machine learning** - detect and extract patterns, adapt to new circumstances
improve ability to generate effective behavior

- **computer vision** - perceive objects
 - **robotics** - manipulate objects
- interact beyond language communication


Machine Learning

How do we define “Learning”?

1. experience E (data)
2. task T
3. performance measure P

An agent learns when its performance at task T , measure by P , improves with experience E .

We need to get a computer do well on a task without manually building its competence.



Machine Learning

How do we define “Learning”?

1. experience E (data)
2. task T
3. performance measure P



T : Playing checkers

E : Playing practice games against itself

P : Percentage of games won against an arbitrary opponent

T : Recognizing hand-written words

E : Database of human-labeled images of handwritten words

P : Percentage of words correctly classified

T : Driving on four-lane highways using vision sensors

E : A sequence of images and steering commands recorded while observing a human driver.


P : Average distance traveled before a human-judged error

T : Categorize email messages as spam or legitimate

E : Database of emails, some with human-given labels

P : Percentage of email messages correctly classified

Programming with Data




We want adaptive, robust and fault tolerant systems

Rule based implementation (often)

- is difficult (for programmers)
- is brittle (can miss many edge cases)
- becomes a nightmare to maintain explicitly
- does not work too well

Programming with Data



We want adaptive, robust and fault tolerant systems


Rule based implementation (often)

- is difficult (for programmers)
- is brittle (can miss many edge cases)
- becomes a nightmare to maintain explicitly
- does not work too well



figure credit: G. Hinton

Programming with Data



We want adaptive, robust and fault tolerant systems
Rule based implementation (often)

- is difficult (for programmers)
- is brittle (can miss many edge cases)
- becomes a nightmare to maintain explicitly
- does not work too well





figure credit: A. Smola

Programming with Data




We want adaptive, robust and fault tolerant systems

Rule based implementation (often)

- is difficult (for programmers)
- is brittle (can miss many edge cases)
- becomes a nightmare to maintain explicitly
- does not work too well



Programming with Data




We want adaptive, robust and fault tolerant systems

Rule based implementation (often)

- is difficult (for programmers)
 - is brittle (can miss many edge cases)
 - becomes a nightmare to maintain explicitly
 - does not work too well
-
- It is easier to obtain examples of what we want: **IF x THEN DO y**
 - collect many pairs (x_i, y_i)
 - estimate a function f such that $f(x_i) = y_i$ (supervised learning)
 - detect patterns in data (unsupervised learning)

Programming with Data




- It is easier to obtain examples of what we want: **IF x THEN DO y**
- collect many pairs (x_i, y_i)
- estimate a function f such that $f(x_i) = y_i$ (supervised learning)
- detect patterns in data (unsupervised learning)

Supervised Learning: Classification

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function $f(x)$ to predict y given x
 - y is categorical == classification

Binary Classification


$y=\{-1,1\}$



Supervised Learning: Classification

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function $f(x)$ to predict y given x
 - y is categorical == classification


Binary Classification $y=\{-1,1\}$



Supervised Learning: Classification

- x can be multi-dimensional
 - Each dimension corresponds to an attribute



Binary Classification
 $y=\{-1,1\}$



- Clump Thickness
- Uniformity of Cell Size
- Uniformity of Cell Shape
- ...

Supervised Learning: Classification

Multiclass Classification
 $y=\{1,...,k\}$




map image x to digit y

Supervised Learning: Regression

- Given $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function $f(x)$ to predict y given x
 - y is real-valued == regression


red: linear

green: non linear




Unsupervised Learning

- Given $x_1, x_2, \dots, x_n$ (without labels)
- Output hidden structure behind the x 's
 - E.g., clustering




Unsupervised Learning

Genomics application: group individuals by genetic similarity



Organize computing clusters





Social network analysis



Market segmentation




Astronomical data analysis



[Variance component model to account for sample structure in genome-wide association studies](#), Nature Genetics 2010

Many more Learnings...

- Reinforcement Learning




Agent and environment interact at discrete time steps : $t = 0, 1, 2, K$

Agent observes state at step t : $s_t \in S$

produces action at step t : $a_t \in A(s_t)$


gets resulting reward : $r_{t+1} \in \mathcal{R}$

and resulting next state : s_{t+1}



Many more Learnings...

- Reinforcement Learning





Agent and environment interact at discrete time steps : $t = 0, 1, 2, K$

Agent observes state at step t : $s_t \in S$

produces action at step t : $a_t \in A(s_t)$


gets resulting reward : $r_{t+1} \in \mathfrak{R}$

and resulting next state : s_{t+1}



Many more Learnings...

- Reinforcement Learning





Agent and environment interact at discrete time steps : $t = 0, 1, 2, K$

Agent observes state at step t : $s_t \in S$

produces action at step t : $a_t \in A(s_t)$

gets resulting reward : $r_{t+1} \in \mathfrak{R}$


and resulting next state : s_{t+1}



New Course
**Robot
Learning**
01HFNOV

Many more Learnings...

- Reinforcement Learning




Agent and environment interact at discrete time steps : $t = 0, 1, 2, K$

Agent observes state at step t : $s_t \in S$

produces action at step t : $a_t \in A(s_t)$

gets resulting reward : $r_{t+1} \in \mathcal{R}$

and resulting next state : s_{t+1}




- Active Learning
- Online Learning
- Multi-task Learning
- Metric Learning
- Ensemble Learning
- Self-Supervised Learning
- Discriminative vs Generative Learning
- [...]

Who is a Machine Learner?




<https://xkcd.com/1838/>

Who is a Machine Learner?




<https://xkcd.com/1838/>

- Framing a Learning Problem
- Designing a Learning System
 - choose the training experience
 - choose exactly what is to be learned (target function)
 - choose a learning algorithm to infer the target function from the experience




Based on a slide of Ray Mooney

Who is a Machine Learner?



<https://xkcd.com/1838/>

- Framing a Learning Problem
- Designing a Learning System
 - choose the training experience
 - choose exactly what is to be learned (target function)
 - choose a learning algorithm to infer the target function from the experience



Based on a slide of Ray Mooney

Pseudo-Algorithm

1. Understand domain, prior knowledge, and goals
2. Data integration, selection, cleaning, pre-processing, etc.
3. Learn models
4. Interpret results
5. Consolidate and deploy discovered knowledge
6. Go back to 1

MORE THAN
243,000 PHOTOS
UPLOADED



MORE THAN
3.8 MILLION
SEARCHES ON
GOOGLE



MORE THAN
350,000
TWEETS
SENT



MORE THAN
65,000
PHOTOS
UPLOADED



MORE THAN
210,000
SNAPS
UPLOADED



120 NEW
ACCOUNTS
CREATED
ON LINKEDIN



MORE THAN
29 MILLION
MESSAGES PROCESSED

1 MILLION PHOTOS

175,000
VIDEO MESSAGES
SHARED



MORE THAN
156 MILLION
E-MAILS SENT



MORE THAN
400 HOURS
OF VIDEOS UPLOADED

70,000
HOURS
OF VIDEO CONTENT
WATCHED

You Tube

AROUND
700,000 HOURS
OF VIDEOS WATCHED



MORE THAN
800,000
FILES
UPLOADED
ON DROPBOX

MORE THAN
5,500 CHECKINS
ON FOURSQUARE



MORE THAN
2,000,000 MINUTES
OF CALLS DONE
BY SKYPE USERS

AROUND
200
EVENT TICKETS
SOLD
ON EVENTBRITE

Eventbrite™

MORE THAN
1000
IMAGES
UPLOADED

imgur

MORE THAN
50 NEW
REVIEWS

MORE THAN
500,000
APPS
DOWNLOADED



MORE THAN
1,000,000
SWIPES

18,000
MATCHES
ON TINDER



16,550 VIDEO
VIEWS
ON VIMEO

MORE THAN
25,000 POSTS
ON TUMBLR

MORE THAN
87,000 HOURS
OF VIDEO
WATCHED

NETFLIX

THINGS THAT
HAPPEN ON INTERNET
EVERY
60
SECONDS


How to deal with data?

- **Observe the data**
 - Summarize their 'uncertainty'
 - Search for constraints and invariances (prior knowledge)
 - Find a model to explain the data
- Represent the data
- Find a solution to learn both the representation and the model for the data by exploiting advanced computing capabilities
- Avoid data bias

How to deal with data?

✗ Unstructured Quartet

Each dataset here also has the same summary statistics. However, they are not *clearly different* or *visually distinct*.




How to deal with data?



Anscombe's Quartet

Each dataset has the same summary statistics (mean, standard deviation, correlation), and the datasets are *clearly different*, and *visually distinct*.















How to deal with data?

Datasaurus dataset

All these datasets
have the same
summary stats to 2
decimal places



X Mean: 54.26
Y Mean: 47.83
X SD : 16.76
Y SD : 26.93
Corr. : -0.06






How to deal with data?

- Observe the data
 - **Summarize their ‘uncertainty’**
 - **Search for constraints and invariances (prior knowledge)**
 - **Find a model to explain the data**
-
- Represent the data
-
- Find a solution to learn both the representation and the model for the data by exploiting advanced computing capabilities
-
- Avoid data bias

Models for describing the data

Main idea: Look at the world, identify what knowledge we have about it, and use this knowledge to construct our model.


The model is a map from input (x) to output (y)




Models for describing the data

Main idea: Look at the world, identify what knowledge we have about it, and use this knowledge to construct our model.


The model is a map from input (x) to output (y)



$$y = ax + b$$



$$y = ax^2 + bx + c$$




$$y = ax^3 + bx^2 + cx + d$$


Models for describing the data

Main idea: Look at the world, identify what knowledge we have about it, and use this knowledge to construct our model.


The model is a map from input (x) to output (y)



$$y = ax + b$$




$$y = ax^2 + bx + c$$



$$y = a \sin(x) + bx + c$$


Choosing a Model Family

The same data can be described in different ways.
What is the “right” way?




Choosing a Model Family

The same data can be described in different ways.
What is the “right” way?




Choosing a Model Family

The same data can be described in different ways.
What is the “right” way?



$$y = ax^2 + bx + c$$

=




$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} a_x \\ a_y \end{pmatrix} t^2 + \begin{pmatrix} b_x \\ b_y \end{pmatrix} t + \begin{pmatrix} c_x \\ c_y \end{pmatrix}$$

Choosing a Model Family


The same data can be described in different ways.
What is the “right” way?

as a **function**



y is not a function of x

as a **parametric curve**




$$\begin{pmatrix} x \\ y \end{pmatrix} = a \begin{pmatrix} \cos t \\ \sin t \end{pmatrix} (a - t)$$

Choosing a Model Family


The same data can be described in different ways.
What is the “right” way?

as a **function**



$r = a\theta$
(polar coordinates)

as a **parametric curve**




$\begin{pmatrix} x \\ y \end{pmatrix} = a \begin{pmatrix} \cos t \\ \sin t \end{pmatrix} (a - t)$

Choosing a Model Family

The same data can be described in different ways.
What is the “right” way?


as a **function**



$r = a\theta$
(polar coordinates)

linear!

as a **parametric curve**




$\begin{pmatrix} x \\ y \end{pmatrix} = a \begin{pmatrix} \cos t \\ \sin t \end{pmatrix} (a - t)$


How to deal with data?

- Observe the data
 - Summarize their 'uncertainty'
 - Search for constraints and invariances (prior knowledge)
 - Find a model to explain the data
-
- **Represent the data**
-
- Find a solution to learn both the representation and the model for the data by exploiting advanced computing capabilities
-
- Avoid data bias

Choosing a Representation




**Representation
Features**




Choosing a Representation



Classification
Model




Representation
Features




Choosing a Representation



Classification
Model




Representation
Features




Choosing a Representation



**Classification
Model**




**Representation
Features**




Choosing a Representation




**Classification
Model**



**Representation
Features**




Choosing a Representation




**Representation
Features**


**Classification
Model**



Choosing a Representation



Linear and Non-linear
Dimensionality
Reduction




How to deal with data?

- Observe the data
 - Summarize their 'uncertainty'
 - Search for constraints and invariances (prior knowledge)
 - Find a model to explain the data
-
- Represent the data
-
- Find a solution to **learn both the representation and the model** for the data by exploiting **advanced computing capabilities**
-
- Avoid data bias


Deep Learning



**Classification
Model**




**Representation
Features**




Deep Learning




**Classification
Model**




**Representation
Features**



Deep Learning



Deep Learning



Deep Learning

Features are Task-driven

Speaking about features only makes sense if we are given a task to solve!




Is color important?

Rank, suit, and color are generic features, but specific problems determine what features are important for that task. What counts in spades, does not count in poker.


Deep Learning

Deep learning is a task-driven paradigm to extract patterns and latent features from given observations. However, features are not the focus of deep learning; rather, they are instrumental for the given task and drive the decision.

Example: Visual classification



Data Bias and Fairness



AI is objective only in the sense of learning what human teaches.

The data provided by human can be highly biased.

Data Bias and Fairness

	WHITE	AFRICAN AMERICAN
Labeled Higher Risk, But Didn't Re-Offend	23.5%	44.9%
Labeled Lower Risk, Yet Did Re-Offend	47.7%	28.0%


Overall, Northpointe's assessment tool correctly predicts recidivism 61 percent of the time. But blacks are almost twice as likely as whites to be labeled a higher risk but not actually re-offend. It makes the opposite mistake among whites: They are much more likely than blacks to be labeled lower risk but go on to commit other crimes. (Source: ProPublica analysis of data from Broward County, Fla.)

AI is objective only in the sense of learning what human teaches.


The data provided by human can be highly biased.

Data Bias - Adaptation needed


From dataset to dataset




From RGB to depth



From simulated to real control




From CAD models to real images



Summary

- Course and Exam presentation
- What is
 - artificial intelligence
 - machine learning
 - deep learning
- Learning: build a model
- How to deal with data
 - observe samples
 - feature representation
 - dataset bias and fairness



<https://xkcd.com/1838/>

How to deal with data

- Observe the data
- **Summarize their ‘uncertainty’**
- **Search for constraints and invariances (prior knowledge)**
- Find a model to explain the data
- Represent the data
- Find a solution to learn both the representation and the model for the data by exploiting advanced computing capabilities
- Avoid data bias

Machine Learner Pseudo-Algorithm

1. **Understand domain, prior knowledge, and goals**
2. Data integration, selection, cleaning, pre-processing, etc.
3. Learn models
4. Interpret results
5. Consolidate and deploy discovered knowledge
6. Go back to 1

Probabilistic reasoning