



Politecnico
di Torino



Data Science and Machine Learning for Engineering Applications

Scikit-learn and Pandas
preprocessing

DataBase and Data Mining Group

Salvatore Greco
Andrea Pasini
Flavio Giobergia
Elena Baralis
Tania Cerquitelli



- **Data preprocessing with Scikit-Learn and Pandas**
- **Summary**
 - **Missing values**
 - **Normalization**
 - **Feature extraction (examples)**
 - Handling nominal data
 - Computing TF-IDF
 - **Dimensionality reduction**
 - PCA



- Scikit-learn estimators are incompatible with **missing values** (e.g., NaN, blank, etc.)
- Scikit-learn estimators assume that all values in an array are **numerical**
- Handling missing values
 - **Discard** entire rows/columns containing missing values
 - **Impute** missing values (e.g., mean, median, constant, etc.)



Discard Missing values



- Create a DataFrame with some **missing values** (NaN values in this case)

```
In [1]: data = np.array([[1, 2, 3, 4],  
                        [5, 6, 7, np.nan],  
                        [9, 10, np.nan, 11]])  
  
df = pd.DataFrame(data, columns=["A", "B", "C", "D"])  
  
df
```

```
Out[1]:
```

	A	B	C	D
0	1.0	2.0	3.0	4.0
1	5.0	6.0	7.0	NaN
2	9.0	10.0	NaN	11.0



Discard Missing values

PoliTo

DBG
M

- You can create a **Mask** with **True** in the presence of **missing values** and False otherwise with *df.isnull()*

In [1]:

```
df.isnull()
```

Out[1]:

	A	B	C	D
0	False	False	False	False
1	False	False	False	True
2	False	False	True	False



Discard Missing values



- You can **count** the number of **missing values** in each column with the `.sum()` method

```
In [1]: df.isnull().sum()
```

```
Out[1]: A    0  
       B    0  
       C    1  
       D    1  
       dtype: int64
```



Discard Missing values



- You can **remove** rows or columns containing missing values with the *df.dropna(axis=)* method
 - Axis=0 to remove rows
 - Axis=1 to remove columns

```
In [1]: df = df.dropna(axis=0) #Remove rows containing missing values
```

```
Out[1]:
```

	A	B	C	D
0	1.0	2.0	3.0	4.0

Note that all rows containing at least one column with a missing value are removed

```
In [1]: df = df.dropna(axis=1) #Remove columns containing missing values
```

```
Out[1]:
```

	A	B
0	1.0	2.0
1.	5.0	6.0
2	9.0	10.0

Note that all columns containing at least one row with a missing value are removed



Discard Missing values



- You can **remove** rows or columns containing missing values with the *df.dropna(axis=)* method
 - Axis=0 to remove rows
 - Axis=1 to remove columns

```
In [1]: df = df.dropna(axis=0) #Remove rows containing missing values
```

```
Out[1]:
```

	A	B	C	D
0	1.0	2.0	3.0	4.0

Note that all rows containing at least one column with a missing value are removed

```
In [1]: df = df.dropna(axis=1) #Remove columns containing missing values
```

```
Out[1]:
```

	A	B
0	1.0	2.0
1.	5.0	6.0
2	9.0	10.0

Note that all columns containing at least one row with a missing value are removed



Discard Missing values



- You can **remove** rows containing missing values in a **specific column** specifying the subset parameter

```
In [1]: df = df.dropna(subset = ["D"])      #Remove rows containing missing values in the 'D' column
```

```
Out[1]:
```

	A	B	C	D
0	1.0	2.0	3.0	4.0
2	9.0	10.0	NaN	11.0

Note that all rows containing a missing value in the 'D' column are removed

- You can also **select** all the rows with **not-null values** in a specific column

```
In [1]: df = df[df['D'].notnull()]
```



Impute Missing values



- Scikit-Learn provides some built-in functions to **impute** values on missing data
 - ***SimpleImputer*** can **replace missing values** using a **descriptive statistic** (e.g., *mean*, *median*, or *most frequent*) along each column, or using a *constant* value.

In [1]:

```
from sklearn.impute import SimpleImputer
imp_mean = SimpleImputer(missing_values=np.nan, strategy='mean')
imp_mean.fit(X_train.values) # Fit only on training data!
X_train = imp_mean.transform(X_train.values)
X_test = imp_mean.transform(X_test.values)
X_train
```

Out[1]:

```
array([[ 1.,  2.,  3.,  4.],
       [ 5.,  6.,  7.,  7.5],
       [ 9., 10.,  5., 11.]])
```

Note that is returned a Numpy array. If you want, you can create a new DataFrame from it



Impute Missing values



- You can specify how to **impute** the value to replace missing values specifying the *strategy* parameter
 - *"mean"*: replace missing values using the **mean along each column**. It can only be used with numeric data.
 - *"median"*: replace missing values using the median along each column. It can only be used with numeric data.
 - *"most_frequent"*: replace missing using **the most frequent value along each column**. It can be used with strings or numeric data. If there is more than one such value, only the smallest is returned.
 - *"constant"*: replace missing values with *fill_value* parameter. It can be used with strings or numeric data.
 - For string or object data types, *fill_value* must be a string. If None, *fill_value* will be 0 when imputing numerical data and "missing_value" for strings or object data types.



Impute Missing values



- You can **impute** values on missing data with Pandas with the *df.fillna()* method
- You should specify the value that will replace the *NaN* values
 - It can be a constant, or a statistic computed with pandas (e.g., *np.mean()*, *np.median()*, etc.)

In [1]:

```
df.fillna(df.mean()) # Fill NaN values with column mean
```

Out[1]:

	A	B	C	D
0	1.0	2.0	3.0	4.0
1	5.0	6.0	7.0	7.5
2	9.0	10.0	5.0	11.0

Note that, in this case, it is returned a DataFrame



- Examples:
 - min-max normalization: MinMaxScaler
 - z-score normalization: StandardScaler

```
In [1]: from sklearn.preprocessing import MinMaxScaler  
        from sklearn.preprocessing import StandardScaler  
  
        minmax_s = MinMaxScaler()  
        zscore_s = StandardScaler()
```



- Applying normalization to training and test set

```
In [1]: X_train = [[0, 10], [0, 20], [2, 10], [2, 20]]
        X_test = [[1, 15]]

        minmax_s.fit(X_train) # NOTE: "learning" on training data only!
        X_train_norm = minmax_s.transform(X_train)
        X_test_norm = minmax_s.transform(X_test) # correct
        X_test_wrong = minmax_s.fit_transform(X_test) # do not fit on test
        print(X_test_norm)
        print(X_test_wrong)
```

```
Out[1]: [[0.5 0.5]]
        [[0, 0]]
```



ALWAYS REMEMBER

PoliTo

DBG
M

1. Fit the transformers only on the training set
2. Transform the training set (use `fit_transform()` if convenient)
3. Only transform the test set — never fit on it!



- Necessary when a datasets presents samples that:
 - Are **not numerical vectors**
 - Example: nominal data, text, images
 - The **model** has a low capacity/can't extract enough knowledge from the row features
 - Example: extraction of polynomial features
 - Example: extraction of season from dates



- Nominal data
 - Two nominal values can only be compared with the equality operator (**cannot be ordered**)
 - For this reason it is **incorrect** to map them to integer features:
 - E.g. 'red', 'green', 'blue' \rightarrow [0, 1, 2]
 - Colors have no ordering
 - The model could infer ordering properties that do not describe correctly our data



- Nominal data
 - One of the simplest solutions is to use **one-hot encoding**:
 - Red $\rightarrow 0, 0, 1$
 - Green $\rightarrow 0, 1, 0$
 - Blue $\rightarrow 1, 0, 0$
 - Pay attention: the **size of the output vector** is linear with the number of distinct values for the attribute
 - Some models (e.g. KNN, clustering) may have problems while working with high dimensional data

Remind: Curse of dimensionality:

The curse of dimensionality refers to the exponential increase in data sparsity and computational complexity as the number of features (dimensions) grows, making learning and generalization harder.



- Nominal data: 1-Hot vectors from dictionaries

```
In [1]: from sklearn.feature_extraction import DictVectorizer
```

```
vect = DictVectorizer(sparse=False, dtype=int)
```

```
Out[1]:
```

```
data = [{ 'model' : 'a', 'price' : 20000},  
        { 'model' : 'b', 'price' : 10000},  
        { 'model' : 'c', 'price' : 8000},  
        { 'model' : 'a', 'price' : 40000},  
        { 'model' : 'c', 'price' : 8500}]
```

```
print(vect.fit_transform(data))
```



- Nominal data: 1-Hot vectors from dictionaries

In [1]:

```
...  
print(vect.fit_transform(data))
```

```
data = [{'model' : 'a', 'price' : 20000},  
        {'model' : 'b', 'price' : 10000},  
        {'model' : 'c', 'price' : 8000},  
        {'model' : 'a', 'price' : 40000},  
        {'model' : 'c', 'price' : 8500}]
```

Out[1]:

```
[[ 1  0  0 20000]  
 [ 0  1  0 10000]  
 [ 0  0  1  8000]  
 [ 1  0  0 40000]  
 [ 0  0  1  8500]]
```

↑
a

↑
b

↑
c



- Nominal data: 1-Hot vectors from dictionaries
 - If you have training and test data use fit and transform separately:

```
In [1]: train = data[:3]
        test = data[3:]

        vect = DictVectorizer(sparse=False, dtype=int)
        vect.fit(train) # Learn vocabulary from training set
        test_transformed = vect.transform(test)
```



- 1-Hot encoding with *OneHotEncoder*
 - Allows passing data in tabular form (“feature matrix”)
 - Numerical values are also encoded – **beware!**
 - Some matrix manipulation is required to only encode categorical features

In [1]:

```
from sklearn.preprocessing import OneHotEncoder

X = np.array([
    ["a", 20000],
    ["b", 10000],
    ["c", 8000],
    ["a", 40000],
    ["c", 8500]
])

ohe = OneHotEncoder(sparse=False, dtype=int)
X_ohe = ohe.fit_transform(X[:, :1])

# same result as DictVectorizer (from before)
X_enc = np.hstack([ X_ohe, X[:, 1:].astype(int)])
```



- Textual data

- Convert textual documents to count vectors

- 1 feature for each word of the vocabulary that count the number of occurrences in the document
- Scikit-learn transformer: *CountVectorizer*
- Example:
 - “My cat. My dog. My cat.”
 - “My dog. My house.”

cat	dog	house	my
2	1	0	3
0	1	1	1



- Textual data
 - Convert textual documents to count vectors
 - Drawback: frequent words have high scores for almost all documents
 - Solution: **TF-IDF** (Term Freq. Inverse Document Freq.)
 - Penalizes words that are common in all documents
 - Boosts words that are frequent in a document, but not in the others



A look into TF-IDF



- **Term Frequency (TF):**
Measures how often a term t appears in a document d .
Inverse Document Frequency (IDF):

$$TF(t, d) = \frac{\# \text{ times term } t \text{ appears in document } d}{\# \text{ total terms in } d}$$

- **Inverse Document Frequency (IDF) :**
Reduces weight for terms common across many documents.

$$IDF(t) = \log \left(\frac{N}{1 + \# \text{ Document containing } t} \right)$$

- $TF\text{-}IDF = TF \times IDF$
High TF-IDF \rightarrow word is frequent in one doc but rare in corpus



■ Textual data: TF-IDF

```
In [1]: from sklearn.feature_extraction.text import TfidfVectorizer  
vect = TfidfVectorizer(stop_words="english")  
  
data = ["dog bites cat", "cat bites dog", "cat and dog house"]  
print(vect.fit_transform(data).toarray())
```

convert to Numpy array

```
Out[1]: [[0.67325467 0.52284231 0.52284231 0.          ]  
[0.67325467 0.52284231 0.52284231 0.          ]  
[0.          0.45329466 0.45329466 0.76749457]]
```



■ Textual data: TF-IDF

```
In [1]: ...  
data = ["dog bites cat", "cat bites dog", "cat and dog house"]  
print(vect.fit_transform(data).toarray())  
# Print the learned vocabulary  
print(vect.vocabulary_)
```

```
Out[1]:
```

	bites	cat	dog	house	
	[0.67325467	0.52284231	0.52284231	0.	Doc 1
	[0.67325467	0.52284231	0.52284231	0.	Doc 2
	[0.	0.45329466	0.45329466	0.76749457]	Doc 3

```
{'dog': 2, 'bites': 0, 'cat': 1, 'house': 3}
```

stopword "and"
has been
removed

specific of this
document



Dimensionality reduction


PoliTo

DBG
MG

- Useful when you want to reduce the number of features for high-dimensional data
 - For **graphical** representations
 - Before applying **classification** and **clustering** to give the features matrix a more **compact** representation



- Example: **PCA**
 - Reduces the dimensionality by finding the directions in the space where data has more variance
 - *If you're into linear algebra:* PCA is just a rotation along the covariance matrix's eigenvectors!





- PCA with Scikit-learn

```
from sklearn.decomposition import PCA

pca = PCA(n_components=5)
X_projection = pca.fit_transform(X)
```

- **n_components** specify the number of components that you want to keep after applying PCA
 - Should be \leq the number of initial features
- The result is a features matrix with the specified number of features



- **Normalize Before PCA!**

PCA is sensitive to the
scale of the features.



Example: think to a screw:

- The **length** is much greater than the **width**
- Without normalization, PCA will align mostly with the length!
- **Fix:** standardize your data (zero mean, unit variance) **before** applying PCA!

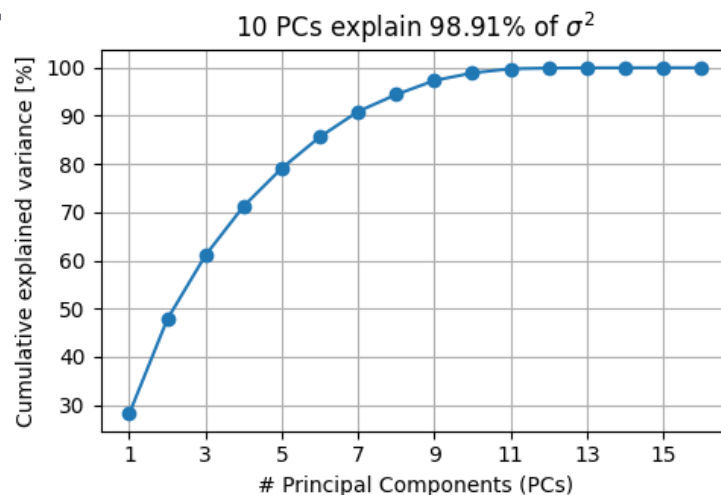


Dimensionality reduction

- Choosing the correct number of components
 - when you do not have too much features:

Cumulative explained variance \geq th.

```
pca = PCA() # compute all the components  
th = 0.9  
X_projection = pca.fit_transform(X)  
cumul_exp_var = np.cumsum(explained_variance)  
n_components = np.where(cum_var >= threshold)[0][0]
```






Dimensionality reduction (when too much features)



- Choosing the correct number of components

```
pca = PCA(n_components=130) #compute the top-N components with largest variance
X_projection = pca.fit_transform(X)
plt.plot(pca.explained_variance_ratio_, marker='o', linestyle='')
```





- Applying the transformation and a classifier

```
pca = PCA(n_components=6)
X_projection = pca.fit_transform(X_train)
my_classifier.train(X_projection, y_train)

# PCA is already fit on training data: do not fit it on test set!
X_test_proj = pca.transform(X_test)
y_test_pred = my_classifier.predict(X_test_proj)
```



- **Other preprocessing methods**
 - <https://scikit-learn.org/stable/modules/classes.html#module-sklearn.preprocessing>